



KubeCon



CloudNativeCon

China 2025





KubeCon



CloudNativeCon

China 2025

# New pattern for sailing multi-host LLM Inference



Kante Yin  
@kerthcet  
<https://ky.dev>  
LWS maintainer, InftyAI farmer




# Why Multi-Host Inference

# Why Multi-Host Inference

- Growing model size

- DeepSeek R1 FP8: 671GB
- KVCache: 400GB
- Needs 16 \* H100 80GB



# Why Multi-Host Inference

- **Growing model size**

- DeepSeek R1 FP8: 671GB
- KVCache: 400GB
- Needs 16 \* H100 80GB

- **High Throughput**

- xPyD paradigm
- Mitigate interference between P & D
- Resource coupling




Figure 6: DistServe Runtime System Architecture


# Why Multi-Host Inference

- Growing model size

- DeepSeek R1 FP8: 671GB
- KVCache: 400GB
- Needs 16 \* H100 80GB

- High Throughput

- xPyD paradigm
- Mitigate interference between P & D
- Resource coupling




from [Perplexity](#)

# What's LWS

LeaderWorkerSet (LWS) is an API for deploying **a group of pods as a unit** of replication, acting like **Statefulset on Statefulset**.

## Design principles:


- Kubernetes capacity reuse
- 1 leader + n workers as a group (superpod), leader as the proxy
- The superpod should behave as an unit, e.g. lifecycle, rolling update.
- Each Pod should have an unique index because we're sharding, that's why we choose Statefulset



# How it Works


## Workflow

- Create a leaderWorkerSet
- Create a leader Statefulset with replicas=4
- Each leader Pod creates a worker Statefulset with replicas= 2.
- The leader Pod and worker Pods are grouped as the superpod



# What it Looks Like


```
apiVersion: leaderworkerset.x-k8s.io/v1
kind: LeaderWorkerSet
metadata:
  name: lws-ample
spec:
  replicas: 4
  leaderWorkerTemplate:
    size: 3
    leaderTemplate:
      spec:
        ...
      workerTemplate:
        spec:
          ...
    rolloutStrategy:
      type: RollingUpdate
      rollingUpdateConfiguration:
        maxUnavailable: 2
        maxSurge: 2
    restartPolicy: RecreateGroupOnPodRestart
```



# Features


# Features - HPA

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: lws-hpa
spec:
  minReplicas: 3
  maxReplicas: 5
  metrics:
    - type: Resource
      resource:
        name: cpu
      target:
        type: Utilization
        averageUtilization: 50
  scaleTargetRef:
    apiVersion: leaderworkerset.x-k8s.io/v1
    kind: LeaderWorkerSet
    name: lws-sample
```




# Feature - TopologyAwarePlacement

```
apiVersion: leaderworkerset.x-k8s.io/v1
kind: LeaderWorkerSet
metadata:
  name: lws-sample
  annotations:
    leaderworkerset.sigs.k8s.io/exclusive-topology: nodepool
spec:
  replicas: 3
  leaderWorkerTemplate:
    size: 9
  ...
  ...
```



# Feature - SubGroup

```
apiVersion: leaderworkerset.x-k8s.io/v1
kind: LeaderWorkerSet
metadata:
  name: lws-sample
spec:
  replicas: 3
  leaderWorkerTemplate:
    size: 7
    subGroupPolicy:
      type: LeaderExcluded
      subGroupSize: 2
  ...
  ...
```



One Replica of LWS (P:D=2:1)


# Ongoing Functions

## KEPs

- Gang scheduling support, [#407](#) — Volcano community
- ResourceClaimTemplate integration, [#444](#)
- Fine-gained system metrics, like rolling update duration, [#87](#)
- Make spec.leaderWorkerTemplate.size mutable, [#552](#) — MistralAI
- Partitioned update support, [#511](#) — Alibaba
- In-place rolling update for image, [#376](#) — OpenKruise community
- Unique node selector and toleration per replica, [#223](#)

*Mostly planned for v0.7.0, see [issues](#).*

# Disaggregated Serving




## Homogeneous Disaggregated Serving with LWS:

- Proxy can't scale independently
- Static PD ratio, scaling PD as a whole
- With identical Pod template, no separate resource pools for P & D

# Disaggregated Serving

## Heterogeneous Disaggregated Serving:


- Multiple role templates, e.g. Proxy, Prefill, Decode
- Rolling update strategy based on P-D ratio
- Independent scaling capacity



# Disaggregated Serving

## Heterogeneous Disaggregated Serving:

- Multiple role templates, e.g. Proxy, Prefill, Decode
- Rolling update strategy based on P-D ratio
- Independent scaling capacity



We're looking to build a new orchestration on top of LWS.

# Adopters & Integrations

## Adopters:

- AWS
- DaoCloud
- Google Cloud
- Nvidia
- More than we know ...

## Integrations:

- Nvidia Dynamo
- Ilmaz
- vLLM
- SGLang

Learn more details at our [website](#). Please join the list if you use LWS as well.

# Join Us



**Github:** <https://github.com/kubernetes-sigs/lws>

**Website:** <https://lws.sigs.k8s.io/>

**Slack:** we're under the guidance of [wg-serving](#)

♥ Thanks to all the contributors!

# Join Us



**Github:** <https://github.com/kubernetes-sigs/lws>

**Website:** <https://lws.sigs.k8s.io/>

**Slack:** we're under the guidance of [wg-serving](#)



Give us a star !

# Join Us



**Github:** <https://github.com/kubernetes-sigs/lws>

**Website:** <https://lws.sigs.k8s.io/>

**Slack:** we're under the guidance of [wg-serving](#)



Join us !



KubeCon



CloudNativeCon

China 2025

Thanks!

